

# Container Orchestration using Swarm

Presented by :  
Suyash Tandon  
Vinen Furtado

# What is container orchestration ?


- Organizing components of applications at network level
- Container Orchestration allows user to control the containers, group them into clusters and coordinate the processes.
- So basically it's a way to manage containers.

# Why is container orchestration needed?

- Containers on a single docker host for development purposes is a easy to use and a straightforward experience.
- But, what about when we need to deploy our applications for usage? In a production environment?
- We use Container orchestration.
- Why?
- Difficult to have the ability to scale our services on a single host.

# About Swarm architecture

- Docker Swarm is a clustering and scheduling tool for Docker containers
- Used to manage a cluster of Docker nodes as a single virtual system
- Swarm consists of manager node(s), worker node(s) and services
- A worker node will only get instructions but the manager node can give instructions



# 3 Parts of our Presentation

- *How Services can be scheduled on Docker Swarm?*
- *How services can be consumed on Docker Swarm?*
- *How updating services works in Docker Swarm? With a small Demo*

# Services in docker swarm can run in 2 modes:

We have used the official NGINX docker image to create the nginx service.

- Default mode
  - have a specific number of replicated services(here 5)
  - Distributed among all the nodes to balance the load and for flexibility

```
ubuntu@ip-172-31-28-20:~$ docker service ps --format 'table {{.ID}}\t{{.Name}}\t{{.Node}}\t{{.CurrentState}}' nginx
```

ID	NAME	NODE	CURRENT STATE
igssmzpg01uy	nginx.1	worker3	Running 9 seconds ago
uj3ebqt7b215	nginx.2	worker2	Running 9 seconds ago
3kuk1hux067q	nginx.3	worker4	Running 9 seconds ago
rp5d130ch9wx	nginx.4	worker1	Running 9 seconds ago
r7h6pthifopu	nginx.5	master	Running 9 seconds ago

- Global mode
  - Schedules single task on every node
  - Good way to keep a track on which node the task is running

```
ubuntu@ip-172-31-28-20:~$ docker service ps --format 'table {{.ID}}\t{{.Name}}\t{{.Node}}\t{{.CurrentState}}' nginx
```

ID	NAME	NODE	CURRENT STATE
pu1ros64aqy2	nginx.sjx781xu7gsohy87rdcxnz78i	worker3	Running 11 seconds ago
j5cjz3yd6k5y	nginx.ls8y6l219kurps71w7svph00w	worker2	Running 11 seconds ago
z0jpvhe6tu5q	nginx.lx0t2ogx7p5da6lq29td8846z	master	Running 11 seconds ago
y1b3qhv25yzq	nginx.oj7m1drsq5jj7qrsq0nkqkl46	worker1	Running 11 seconds ago
oajtf1xjmkms	nginx.wq6p8avd7da8lrzmzk6l437a	worker4	Running 11 seconds ago

```
ubuntu@ip-172-31-28-20:~$ docker node ls
```

ID	HOSTNAME	STATUS	AVAILABILITY	MANAGER STATUS	ENGINE VERSION
lx0t2ogx7p5da6lq29td8846z *	master	Ready	Active	Leader	18.09.0
oj7m1drsq5jj7qrsq0nkqkl46	worker1	Ready	Active		18.09.0
ls8y6l219kurps71w7svph00w	worker2	Ready	Active		18.09.0
sjx781xu7gsohy87rdcxnz78i	worker3	Ready	Active		18.09.0
wq6p8avd7da8lrzmzk6l437a	worker4	Ready	Active		18.09.0



# Scheduling services

## Strategies

- Spread
  - Its a default strategy where in the service tasks are scheduled on the available node
  - Ensures that the tasks are evenly spread over the nodes and clusters
  - if a new node is created and if there are any tasks yet to be scheduled, then those tasks are scheduled on that node.

```
ubuntu@ip-172-31-28-20:~$ docker service ps --format 'table {{.ID}}\t{{.Name}}\t{{.Node}}\t{{.CurrentState}}' nginx-01
```

ID	NAME	NODE	CURRENT STATE
grbs7hw43ap7	nginx-01.1	worker2	Running 12 seconds ago
qurr8rwoxfat	nginx-01.2	worker1	Running 12 seconds ago
r49ixmd3oh1d	nginx-01.3	worker3	Running 12 seconds ago

```
ubuntu@ip-172-31-28-20:~$ docker service create --name nginx-02 nginx
u37ag0d3kajgc8tyzy0j93qto
overall progress: 1 out of 1 tasks
1/1: running [=====>]
verify: Service converged
ubuntu@ip-172-31-28-20:~$ docker service ls
```

ID	NAME	MODE	REPLICAS	IMAGE	PORTS
wv9a1xus5y2s	nginx-01	replicated	3/3	nginx:latest	
u37ag0d3kajg	nginx-02	replicated	1/1	nginx:latest	

```
ubuntu@ip-172-31-28-20:~$ docker service ps --format 'table {{.ID}}\t{{.Name}}\t{{.Node}}\t{{.CurrentState}}' nginx-02
```

ID	NAME	NODE	CURRENT STATE
2g0rkwqmtvx	nginx-02.1	master	Running 23 seconds ago

- Specific node
  - Tasks can be scheduled on the node or cluster specified.

# Reserve Resources

If the reserved memory of the scheduled task is greater than the memory of any node, that task will not be scheduled on that node.

It is automatically scheduled on a different node.

```
ubuntu@ip-172-31-28-20:~$ docker service create --name nginx-01 --reserve-memory 900Mb --replicas 3 nginx
09e2h1my0p5aj17083js61ogt
overall progress: 3 out of 3 tasks
1/3: running [=====>]
2/3: running [=====>]
3/3: running [=====>]
verify: Service converged
ubuntu@ip-172-31-28-20:~$ docker service ps --format 'table {{.ID}}\t{{.Name}}\t{{.Node}}\t{{.CurrentState}}' nginx-01


| ID           | NAME       | NODE    | CURRENT STATE          |
|--------------|------------|---------|------------------------|
| lu8pneakqlfc | nginx-01.1 | master  | Running 21 seconds ago |
| vbo678khboxh | nginx-01.2 | worker2 | Running 21 seconds ago |
| 8nqx150pzdqp | nginx-01.3 | worker3 | Running 21 seconds ago |


ubuntu@ip-172-31-28-20:~$ docker service create --name nginx-02 --reserve-memory 200Mb --replicas 4 nginx
qby5mfq30fkr1fhdtvo4plf0m
overall progress: 4 out of 4 tasks
1/4: running [=====>]
2/4: running [=====>]
3/4: running [=====>]
4/4: running [=====>]
verify: Service converged
ubuntu@ip-172-31-28-20:~$ docker service ps --format 'table {{.ID}}\t{{.Name}}\t{{.Node}}\t{{.CurrentState}}' nginx-01


| ID           | NAME       | NODE    | CURRENT STATE          |
|--------------|------------|---------|------------------------|
| lu8pneakqlfc | nginx-01.1 | master  | Running 50 seconds ago |
| vbo678khboxh | nginx-01.2 | worker2 | Running 50 seconds ago |
| 8nqx150pzdqp | nginx-01.3 | worker3 | Running 50 seconds ago |


ubuntu@ip-172-31-28-20:~$ docker service ps --format 'table {{.ID}}\t{{.Name}}\t{{.Node}}\t{{.CurrentState}}' nginx-02


| ID           | NAME       | NODE    | CURRENT STATE          |
|--------------|------------|---------|------------------------|
| rews5eqvs6ty | nginx-02.1 | worker1 | Running 33 seconds ago |
| hth4mrs5wlg3 | nginx-02.2 | worker4 | Running 33 seconds ago |
| pkxcs5ac0tu1 | nginx-02.3 | worker1 | Running 33 seconds ago |
| e44adh0nkakw | nginx-02.4 | worker4 | Running 33 seconds ago |


```

- Zone based
  - In this the nodes can be clustered which can be called a zone
  - Scheduling services only on that zone is zone based scheduling.
  - If there are zones created and spread scheduling is used, then the zone will be considered as a single unit node
  - The service tasks scheduled on that zone will be distributed evenly among the nodes on the zone.

```
ubuntu@ip-172-31-28-20:~$ docker node update --label-add 'com.acme.zone=a' master
master
ubuntu@ip-172-31-28-20:~$ docker node update --label-add 'com.acme.zone=a' worker1
worker1
ubuntu@ip-172-31-28-20:~$ docker node update --label-add 'com.acme.zone=b' worker2
worker2
ubuntu@ip-172-31-28-20:~$ docker node update --label-add 'com.acme.zone=c' worker3
worker3
ubuntu@ip-172-31-28-20:~$ docker node update --label-add 'com.acme.zone=d' worker4
worker4
```

```
ubuntu@ip-172-31-28-20:~$ docker service ps --format 'table {{.ID}}\t{{.Name}}\t{{.Node}}\t{{.CurrentState}}' nginx
```

ID	NAME	NODE	CURRENT STATE
4fg17bq1ibqx	nginx.1	worker2	Running 10 seconds ago
133qbxfc5uyd	nginx.2	worker4	Running 10 seconds ago
vk7areq0tmt	nginx.3	worker4	Running 10 seconds ago
p32jicg0i5ei	nginx.4	worker4	Running 10 seconds ago
7ig4qu0tr56i	nginx.5	worker2	Running 10 seconds ago
e7i03x6frnvi	nginx.6	worker2	Running 10 seconds ago
sjs758tx8ulh	nginx.7	master	Running 10 seconds ago
uuhr5kkv1td	nginx.8	worker1	Running 10 seconds ago
8pv3oxppb37m	nginx.9	worker3	Running 10 seconds ago
sf0d2gszc8xw	nginx.10	worker1	Running 10 seconds ago
qfy03li0tz5u	nginx.11	worker4	Running 10 seconds ago
y9vzaqtm707l	nginx.12	worker3	Running 10 seconds ago
3kic64rw52b0	nginx.13	worker3	Running 10 seconds ago
q9f8r193muc6	nginx.14	worker3	Running 10 seconds ago
mx8qj52nrcrg	nginx.15	master	Running 10 seconds ago
d6v8c4h3el0l	nginx.16	worker2	Running 10 seconds ago

- Rescheduling on failure
  - If a node fails, the state of the tasks scheduled on that node and also the current state of the running process is maintained and is scheduled on another node.

```
ubuntu@ip-172-31-28-20:~$ docker service ps --format 'table {{.ID}}\t{{.Name}}\t{{.Node}}\t{{.CurrentState}}' nginx
```

ID	NAME	NODE	CURRENT STATE
kg9d5lv92awg	nginx.1	master	Running 23 seconds ago
zdkhw4pd5u81	nginx.2	worker4	Running 23 seconds ago
yovg123zuhct	nginx.3	worker2	Running 23 seconds ago
0qtlcivddplo	nginx.4	worker3	Running 23 seconds ago
rbz43xodi8ts	nginx.5	master	Running 7 seconds ago
tn4z42jjm0qt	\_ nginx.5	worker1	Shutdown 8 seconds ago

```
ubuntu@ip-172-31-28-20:~$ docker node ls
```

ID	HOSTNAME	STATUS	AVAILABILITY	MANAGER STATUS	ENGINE VERSION
lx0t2ogx7p5da6lq29td8846z *	master	Ready	Active	Leader	18.09.0
oj7m1drsq5jj7qrsq0nkqkl46	worker1	Ready	Drain		18.09.0
ls8y6l219kurps71w7svph00w	worker2	Ready	Active		18.09.0
sjx781xu7gsohy87rdcxnz78i	worker3	Ready	Active		18.09.0
wq6p8avd7da8lrmzmkh6l437a	worker4	Ready	Active		18.09.0

# Consuming services

- Internal service consumption
  - Virtual ip
    - When a new service is created, each service is given a virtual IP address through which it can be accessed,
  - Embedded DNS server
    - Remembering the IP addresses is difficult, as a result swarm makes use of embedded DNS server
    - A query is fired from the service to the DNS server



# Load balancing

- External service consumption
  - Routing mesh
    - A service's port is published on all the nodes in the cluster.
    - Whenever there is an external request, the request is redirected to any free node that has the service's port, even if the node is not running a service task.
    - The swarm then makes use of the routing mesh to redirect the request to the desired node.
  - Host mode
    - The task's port is mapped directly on the node's port and all the service request are redirected on that port.

# Updating Services

**3 modes:**

- **Service Updates**
- **Image Updates**
- **Rollback Updates**

# Service Updates

Using the **docker service update --help** command, we can update just about any aspect of our services.

We tried updating the ENV variable of the service. Added the NODE ID.

```
ubuntu@ip-172-31-28-20:~/src/dockercoins$ docker container inspect --format '{{json .Config.Env}}' $(docker container ls -lq) | jq '.'  
[  
  "PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin"  
]
```

After updating:

```
ubuntu@ip-172-31-28-20:~/src/dockercoins$ docker container inspect --format '{{json .Config.Env}}' $(docker container ls -lq) | jq '.'  
[  
  "NODE_ID=sjx781xu7gsohy87rdcxnz78t",  
  "PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin"  
]
```

# Updating Images

One of the more obvious service updates that might be required is to update the image that the service is based on.

```
ubuntu@ip-172-31-28-20:~/src/dockercoins$ docker service create --detach=false --name nginxhello --publish published=8080,target=80 --mode global nbrown/nginxhello:1.12.1
vbfyr78itgqk9jsr0ugh23kuo
overall progress: 5 out of 5 tasks
wq6p8avd7da8: running [=====>]
lx0t2ogx7p5d: running [=====>]
ls8y6l219kur: running [=====>]
sjx781xu7gso: running [=====>]
oj7m1drsq5jj: running [=====>]
verify: Service converged
ubuntu@ip-172-31-28-20:~/src/dockercoins$ docker service inspect --format '{{.Spec.TaskTemplate.ContainerSpec.Image}}' nginxhello
nbrown/nginxhello:1.12.1@sha256:7b6aab518939e8dff1224e74f6e1f58381b0c4a110f465c4fffdcb69a885a824
ubuntu@ip-172-31-28-20:~/src/dockercoins$ (docker service update --quiet --detach=true --image $IMAGE \
> > --update-delay 15s --update-order "start-first" nginxhello &) > /dev/null 2>&1 && \
> > watch -n 0.1 "docker service ps --filter "desired-state=running" --format '$FORMAT' nginxhello"
-n: command not found
ubuntu@ip-172-31-28-20:~/src/dockercoins$ (docker service update --quiet --detach=true --image $IMAGE > --update-delay 15s --update-order "start-first" nginxhello &) > /dev
/null 2>&1 && > watch -n 0.1 "docker service ps --filter "desired-state=running" --format '$FORMAT' nginxhello"
-n: command not found
ubuntu@ip-172-31-28-20:~/src/dockercoins$ (docker service update --quiet --detach=true --image $IMAGE nginxhello &) > /dev/null 2>&1 && \
> watch -n 1 "docker service ps --filter "desired-state=running" --format '$FORMAT' nginxhello"
ubuntu@ip-172-31-28-20:~/src/dockercoins$ docker service inspect --format '{{.Spec.TaskTemplate.ContainerSpec.Image}}' nginxhello
nbrown/nginxhello:1.13.5@sha256:c0f272bd3b059efa32125abb061aa94935737e87ddbcd2015fa9f2feddd3d9bc
```

# Controlled Service Updates

We might, for example, want to update all of the tasks simultaneously, or two at a time, rather than the default behaviour of one at a time.

We can enforce a rollback when an update fails , instead of pause the update, which is the default behaviour.

**DEMO**

# Rolling Back the Updates

If we hit a problem with an update, we can automatically rollback the update, based on the 'update failure action' policy.

We have simulated an update failure, by simply stopping one of the containers that make up the service's tasks.

DEMO

# What we learnt?

- How service scheduling works and various ways in which a services are scheduled.
- Understand load balancing and how its performed.
- Different ways of updating the already existing services.

Thank you!